

TEMPORAL CREDIT WITHOUT REWARDS: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot skill learning. The motivating bet is: Assign physical credit across long horizons without reward labels or RL framing.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For robot skill learning, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a; 2025; 2023; 2026b;c;d). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.550	0.450
observed_only_wam	hidden_mode_shift	0.205	0.795
augmentation_as_data_wam	nominal	0.603	0.397
augmentation_as_data_wam	hidden_mode_shift	0.260	0.740
scalar_uncertainty_wam	nominal	0.670	0.330
scalar_uncertainty_wam	hidden_mode_shift	0.290	0.710
proposed_branch_mechanism	nominal	0.772	0.228
proposed_branch_mechanism	hidden_mode_shift	0.522	0.478

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Temporal Credit Without Rewards argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Clip-reid: Exploiting vision-language model for image re-identification without concrete text labels, 2023. <https://doi.org/10.1609/aaai.v37i1.25225>.
- Compliant motion planning integrating human skill for robotic arm collecting tomato bunch based on improved ddpg, 2025. <https://doi.org/10.3390/plants14050634>.
- Recovla: Vlm-guided reward compilation for failure recovery in vision-language-action policies, 2026a. <http://arxiv.org/abs/2606.09630v1>.
- Pi-vla: Adaptive symmetry-aware decision-making for long-horizon vision-language-action manipulation, 2026b. <https://doi.org/10.3390/sym18030394>.
- Forecasting long-term spatial-temporal dynamics with generative transformer networks, 2026c. <https://doi.org/10.31224/4295>.
- Evaluating long-range temporal structure in foundation model-based forecasts of heartbeat dynamics, 2026d. <https://doi.org/10.64898/2026.05.25.727760>.